

ISYE 6501 HW Wk 3 Solution

Question 7.1

Describe a situation or problem from your job, everyday life, current events, etc., for which exponential smoothing would be appropriate. What data would you need? Would you expect the value of alpha (the first smoothing parameter) to be closer to 0 or 1, and why?

Working in the supply chain industry, we use exponential smoothing to estimate the package volume level that arrives to the processing terminals. We will need the historical weekly volume at each terminal, with seasonal trend across the year. With the estimated volume level, will allocate the resource to meet the capacity accordingly.

The parameter should be closer to 1 because the volume on weekly level is considered stable with low level randomness.

Question 7.2

Using the 20 years of daily high temperature data for Atlanta (July through October) from Question 6.2 (file temps.txt), build and use an exponential smoothing model to help make a judgment of whether the unofficial end of summer has gotten later over the 20 years. (Part of the point of this assignment is for you to think about how you might use exponential smoothing to answer this question. Feel free to combine it with other models if you'd like to. There's certainly more than one reasonable approach.)

Note: in R, you can use either `HoltWinters` (simpler to use) or the `smooth` package's `es` function (harder to use, but more general). If you use `es`, the Holt-Winters model uses `model="AAM"` in the function call (the first and second constants are used "A"dditively, and the third (seasonality) is used "M"ultiplicatively; the documentation doesn't make that clear).

```
# read the file temps.txt into R
df_raw <- read.table("Wk3/temps.txt", header=TRUE)

# preview and explore df_raw

head(df_raw)
```

##	DAY	X1996	X1997	X1998	X1999	X2000	X2001	X2002	X2003	X2004	X2005	X2006
## 1	1-Jul	98	86	91	84	89	84	90	73	82	91	93
## 2	2-Jul	97	90	88	82	91	87	90	81	81	89	93
## 3	3-Jul	97	93	91	87	93	87	87	87	86	86	93
## 4	4-Jul	90	91	91	88	95	84	89	86	88	86	91
## 5	5-Jul	89	84	91	90	96	86	93	80	90	89	90
## 6	6-Jul	93	84	89	91	96	87	93	84	90	82	81
##	X2007	X2008	X2009	X2010	X2011	X2012	X2013	X2014	X2015			
## 1	95	85	95	87	92	105	82	90	85			
## 2	85	87	90	84	94	93	85	93	87			
## 3	82	91	89	83	95	99	76	87	79			
## 4	86	90	91	85	92	98	77	84	85			
## 5	88	88	80	88	90	100	83	86	84			
## 6	87	82	87	89	90	98	83	87	84			

1. To prepare the input data for exponential modeling

```

# remove the first column to use the data from 1996 to 2015 as input
df_input <- df_raw[,2:21]

# by checking function HoltWinters, the data input x should be an object of class ts - to convert input
# by checking function ts, the data should be a vector or matrix - so to convert the input data first to vector

# first to convert the dataframe to 1-dimension to connect the data across all years
df_input <- unlist(df_input)

# convert 1-dimensional array to vector
df_input_1 <- as.vector(df_input)

# convert to ts, data from 1996 with 123 obs
data_input <- ts(df_input_1, start=1996, frequency=123)

# structure of model input data
str(data_input)

```

```
## Time-Series [1:2460] from 1996 to 2016: 98 97 97 90 89 93 93 91 93 93 ...
```

2. To build model

2.1 To start from single exponential smoothing

```

model_single <- HoltWinters(data_input, beta=FALSE, gamma=FALSE)
model_single

```

```

## Holt-Winters exponential smoothing without trend and without seasonal component.
##
## Call:
## HoltWinters(x = data_input, beta = FALSE, gamma = FALSE)
##
## Smoothing parameters:
##   alpha: 0.8388021
##   beta  : FALSE
##   gamma : FALSE
##
## Coefficients:
##      [,1]
## a 63.30952

```

From single exponential smoothing model, the best alpha is 0.8388021. The Coefficients is 63.30952.

2.2 To build double exponential smoothing model with trend

```

model_double <- HoltWinters(data_input, gamma=FALSE)
model_double

```

```

## Holt-Winters exponential smoothing with trend and without seasonal component.
##
## Call:
## HoltWinters(x = data_input, gamma = FALSE)

```

```
##
## Smoothing parameters:
## alpha: 0.8445729
## beta : 0.003720884
## gamma: FALSE
##
## Coefficients:
##      [,1]
## a 63.2530022
## b -0.0729933
```

From the above model result, the beta is ~ 0.0037 and the Coefficients of beta is ~ -0.07 which are all close to 0. The beta itself is close to 0 that means the trend is more close to randomness instead of actual trend. And the Coefficients b is also close to zero meaning the trend is not significant that we can ignore.

2.3 To build Holt-Winter model (with trend and cyclic pattern)

```
model_triple <- HoltWinters(data_input)
model_triple
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = data_input)
##
## Smoothing parameters:
## alpha: 0.6610618
## beta : 0
## gamma: 0.6248076
##
## Coefficients:
##      [,1]
## a    71.477236414
## b   -0.004362918
## s1   18.590169842
## s2   17.803098732
## s3   12.204442890
## s4   13.233948865
## s5   12.957258705
## s6   11.525341233
## s7   10.854441534
## s8   10.199632666
## s9    8.694767348
## s10   5.983076192
## s11   3.123493477
## s12   4.698228193
## s13   2.730023168
## s14   2.995935818
## s15   1.714600919
## s16   2.486701224
## s17   6.382595268
## s18   5.081837636
## s19   7.571432660
## s20   6.165047647
```

##	s21	9.560458487
##	s22	9.700133847
##	s23	8.808383245
##	s24	8.505505527
##	s25	7.406809208
##	s26	6.839204571
##	s27	6.368261304
##	s28	6.382080380
##	s29	4.552058253
##	s30	6.877476437
##	s31	4.823330209
##	s32	4.931885957
##	s33	7.109879628
##	s34	6.178469084
##	s35	4.886891317
##	s36	3.890547248
##	s37	2.148316257
##	s38	2.524866001
##	s39	3.008098232
##	s40	3.041663870
##	s41	2.251741386
##	s42	0.101091985
##	s43	-0.123337548
##	s44	-1.445675315
##	s45	-1.802768181
##	s46	-2.192036338
##	s47	-0.180954242
##	s48	1.538987281
##	s49	5.075394760
##	s50	6.740978049
##	s51	7.737089782
##	s52	8.579515859
##	s53	8.408834158
##	s54	4.704976718
##	s55	1.827215229
##	s56	-1.275747384
##	s57	1.389899699
##	s58	1.376842871
##	s59	0.509553410
##	s60	1.886439429
##	s61	-0.806454923
##	s62	5.221873550
##	s63	5.383073482
##	s64	4.265584552
##	s65	3.841481452
##	s66	-0.231239928
##	s67	0.542761270
##	s68	0.780131779
##	s69	1.096690727
##	s70	0.690525998
##	s71	2.301303414
##	s72	2.965913580
##	s73	4.393732595
##	s74	2.744547070

```

## s75      1.035278911
## s76      1.170709479
## s77      2.796838283
## s78      2.000312540
## s79      0.007337449
## s80     -1.203916069
## s81      0.352397232
## s82      0.675108103
## s83     -3.169643942
## s84     -1.913321175
## s85     -1.647780450
## s86     -5.281261301
## s87     -5.126493027
## s88     -2.637666754
## s89     -2.342133004
## s90     -3.281910970
## s91     -4.242033198
## s92     -2.596010530
## s93     -7.821281290
## s94     -8.814741200
## s95     -8.996689798
## s96     -7.835655534
## s97     -5.749139155
## s98     -5.196182693
## s99     -8.623793296
## s100    -11.809355220
## s101    -13.129428554
## s102    -16.095143067
## s103    -15.125436350
## s104    -13.963606549
## s105    -12.953304848
## s106    -16.097179844
## s107    -15.489223470
## s108    -13.680122300
## s109    -11.921434142
## s110    -12.035411347
## s111    -12.837047727
## s112     -9.095808127
## s113     -5.433029341
## s114     -6.800835107
## s115     -8.413639598
## s116    -10.912409484
## s117    -13.553826535
## s118    -10.652543677
## s119    -12.627298331
## s120     -9.906981556
## s121    -12.668519900
## s122     -9.805502547
## s123     -7.775306633

```

From the output, the beta and Coefficients b are still close to 0 meaning there is no clear trend of change across 20 years.

2.4 With Seasonal Trend

```
model_season <- HoltWinters(data_input,seasonal="multiplicative")
model_season
```

```
## Holt-Winters exponential smoothing with trend and multiplicative seasonal component.
##
## Call:
## HoltWinters(x = data_input, seasonal = "multiplicative")
##
## Smoothing parameters:
##   alpha: 0.615003
##   beta : 0
##   gamma: 0.5495256
##
## Coefficients:
##               [,1]
## a      73.679517064
## b     -0.004362918
## s1      1.239022317
## s2      1.234344062
## s3      1.159509551
## s4      1.175247483
## s5      1.171344196
## s6      1.151038408
## s7      1.139383104
## s8      1.130484528
## s9      1.110487514
## s10     1.076242879
## s11     1.041044609
## s12     1.058139281
## s13     1.032496529
## s14     1.036257448
## s15     1.019348815
## s16     1.026754142
## s17     1.071170378
## s18     1.054819556
## s19     1.084397734
## s20     1.064605879
## s21     1.109827336
## s22     1.112670130
## s23     1.103970506
## s24     1.102771209
## s25     1.091264692
## s26     1.084518342
## s27     1.077914660
## s28     1.077696145
## s29     1.053788854
## s30     1.079454300
## s31     1.053481186
## s32     1.054023885
## s33     1.078221405
## s34     1.070145761
## s35     1.054891375
## s36     1.044587771
```

s37 1.023285461
s38 1.025836722
s39 1.031075732
s40 1.031419152
s41 1.021827552
s42 0.998177248
s43 0.996049257
s44 0.981570825
s45 0.976510542
s46 0.967977608
s47 0.985788411
s48 1.004748195
s49 1.050965934
s50 1.072515008
s51 1.086532279
s52 1.098357400
s53 1.097158461
s54 1.054827180
s55 1.022866587
s56 0.987259326
s57 1.016923524
s58 1.016604903
s59 1.004320951
s60 1.019102781
s61 0.983848662
s62 1.055888360
s63 1.056122844
s64 1.043478958
s65 1.039475693
s66 0.991019224
s67 1.001437488
s68 1.002221759
s69 1.003949213
s70 0.999566344
s71 1.018636837
s72 1.026490773
s73 1.042507768
s74 1.022500795
s75 1.002503740
s76 1.004560984
s77 1.025536556
s78 1.015357769
s79 0.992176558
s80 0.979377825
s81 0.998058079
s82 1.002553395
s83 0.955429116
s84 0.970970220
s85 0.975543504
s86 0.931515830
s87 0.926764603
s88 0.958565273
s89 0.963250387
s90 0.951644060

```
## s91    0.937362688
## s92    0.954257999
## s93    0.892485444
## s94    0.879537700
## s95    0.879946892
## s96    0.890633648
## s97    0.917134959
## s98    0.925991769
## s99    0.884247686
## s100   0.846648167
## s101   0.833696369
## s102   0.800001437
## s103   0.807934782
## s104   0.819343668
## s105   0.828571029
## s106   0.795608740
## s107   0.796609993
## s108   0.815503509
## s109   0.830111282
## s110   0.829086181
## s111   0.818367239
## s112   0.863958784
## s113   0.912057203
## s114   0.898308248
## s115   0.878723779
## s116   0.848971946
## s117   0.813891909
## s118   0.846821392
## s119   0.819121827
## s120   0.851036184
## s121   0.820416491
## s122   0.851581233
## s123   0.874038407
```

From all 4 exponential models, there is no significant trend across 20 years. In conclusion, we can not prove that the unofficial summer ends is getting later across years.

Question 8.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a linear regression model would be appropriate. List some (up to 5) predictors that you might use.

Use linear regression to build the cost model of logistics. To estimated cost, will need parameter like the weight, the requested time of labor to process for different destination, the distance of delivery, and based rate of vendor. Use historical data to find the best parameter and then use the model to estimate the future cost.

Question 8.2

Using crime data from <http://www.statsci.org/data/general/uscrime.txt> (file `uscrime.txt`, description at <http://www.statsci.org/data/general/uscrime.html>), use regression (a useful R function is `lm` or `glm`) to predict the observed crime rate in a city with the following data: M = 14.0 So = 0 Ed = 10.0 Po1 = 12.0 Po2 = 15.5 LF = 0.640 M.F = 94.0 Pop = 150 NW = 1.1 U1 = 0.120 U2 = 3.6 Wealth = 3200 Ineq = 20.1

Prob = 0.04 Time = 39.0 Show your model (factors used and their coefficients), the software output, and the quality of fit.

Note that because there are only 47 data points and 15 predictors, you'll probably notice some overfitting. We'll see ways of dealing with this sort of problem later in the course.

```
# Read us crime data file into R dataframe
df_crime <- read.table("Wk3/uscrime.txt", header=TRUE)

# Data Explore
head(df_crime)
```

```
##      M So   Ed Po1 Po2   LF   M.F Pop   NW   U1 U2 Wealth Ineq
## 1 15.1  1  9.1  5.8  5.6 0.510 95.0 33 30.1 0.108 4.1 3940 26.1
## 2 14.3  0 11.3 10.3  9.5 0.583 101.2 13 10.2 0.096 3.6 5570 19.4
## 3 14.2  1  8.9  4.5  4.4 0.533 96.9 18 21.9 0.094 3.3 3180 25.0
## 4 13.6  0 12.1 14.9 14.1 0.577 99.4 157 8.0 0.102 3.9 6730 16.7
## 5 14.1  0 12.1 10.9 10.1 0.591 98.5 18 3.0 0.091 2.0 5780 17.4
## 6 12.1  0 11.0 11.8 11.5 0.547 96.4 25 4.4 0.084 2.9 6890 12.6
##      Prob   Time Crime
## 1 0.084602 26.2011   791
## 2 0.029599 25.2999  1635
## 3 0.083401 24.3006   578
## 4 0.015801 29.9012  1969
## 5 0.041399 21.2998  1234
## 6 0.034201 20.9995   682
```

```
str(df_crime)
```

```
## 'data.frame':   47 obs. of  16 variables:
## $ M      : num  15.1 14.3 14.2 13.6 14.1 12.1 12.7 13.1 15.7 14 ...
## $ So      : int   1 0 1 0 1 0 0 1 1 1 0 ...
## $ Ed      : num   9.1 11.3 8.9 12.1 12.1 11 11.1 10.9 9 11.8 ...
## $ Po1     : num   5.8 10.3 4.5 14.9 10.9 11.8 8.2 11.5 6.5 7.1 ...
## $ Po2     : num   5.6 9.5 4.4 14.1 10.1 11.5 7.9 10.9 6.2 6.8 ...
## $ LF      : num   0.51 0.583 0.533 0.577 0.591 0.547 0.519 0.542 0.553 0.632 ...
## $ M.F     : num   95 101.2 96.9 99.4 98.5 ...
## $ Pop     : int   33 13 18 157 18 25 4 50 39 7 ...
## $ NW      : num   30.1 10.2 21.9 8 3 4.4 13.9 17.9 28.6 1.5 ...
## $ U1      : num   0.108 0.096 0.094 0.102 0.091 0.084 0.097 0.079 0.081 0.1 ...
## $ U2      : num   4.1 3.6 3.3 3.9 2 2.9 3.8 3.5 2.8 2.4 ...
## $ Wealth  : int  3940 5570 3180 6730 5780 6890 6200 4720 4210 5260 ...
## $ Ineq    : num   26.1 19.4 25 16.7 17.4 12.6 16.8 20.6 23.9 17.4 ...
## $ Prob    : num   0.0846 0.0296 0.0834 0.0158 0.0414 ...
## $ Time    : num   26.2 25.3 24.3 29.9 21.3 ...
## $ Crime   : int   791 1635 578 1969 1234 682 963 1555 856 705 ...
```

```
summary(df_crime)
```

```
##      M              So              Ed              Po1
## Min.    :11.90   Min.    :0.0000   Min.    : 8.70   Min.    : 4.50
## 1st Qu.:13.00   1st Qu.:0.0000   1st Qu.: 9.75   1st Qu.: 6.25
## Median :13.60   Median :0.0000   Median :10.80   Median : 7.80
```

```
## Mean :13.86 Mean :0.3404 Mean :10.56 Mean : 8.50
## 3rd Qu.:14.60 3rd Qu.:1.0000 3rd Qu.:11.45 3rd Qu.:10.45
## Max. :17.70 Max. :1.0000 Max. :12.20 Max. :16.60
## Po2 LF M.F Pop
## Min. : 4.100 Min. :0.4800 Min. : 93.40 Min. : 3.00
## 1st Qu.: 5.850 1st Qu.:0.5305 1st Qu.: 96.45 1st Qu.: 10.00
## Median : 7.300 Median :0.5600 Median : 97.70 Median : 25.00
## Mean : 8.023 Mean :0.5612 Mean : 98.30 Mean : 36.62
## 3rd Qu.: 9.700 3rd Qu.:0.5930 3rd Qu.: 99.20 3rd Qu.: 41.50
## Max. :15.700 Max. :0.6410 Max. :107.10 Max. :168.00
## NW U1 U2 Wealth
## Min. : 0.20 Min. :0.07000 Min. :2.000 Min. :2880
## 1st Qu.: 2.40 1st Qu.:0.08050 1st Qu.:2.750 1st Qu.:4595
## Median : 7.60 Median :0.09200 Median :3.400 Median :5370
## Mean :10.11 Mean :0.09547 Mean :3.398 Mean :5254
## 3rd Qu.:13.25 3rd Qu.:0.10400 3rd Qu.:3.850 3rd Qu.:5915
## Max. :42.30 Max. :0.14200 Max. :5.800 Max. :6890
## Ineq Prob Time Crime
## Min. :12.60 Min. :0.00690 Min. :12.20 Min. : 342.0
## 1st Qu.:16.55 1st Qu.:0.03270 1st Qu.:21.60 1st Qu.: 658.5
## Median :17.60 Median :0.04210 Median :25.80 Median : 831.0
## Mean :19.40 Mean :0.04709 Mean :26.60 Mean : 905.1
## 3rd Qu.:22.75 3rd Qu.:0.05445 3rd Qu.:30.45 3rd Qu.:1057.5
## Max. :27.60 Max. :0.11980 Max. :44.00 Max. :1993.0
```

```
# Return the column names of data frame in order to build lm model next
colnames(df_crime)
```

```
## [1] "M" "So" "Ed" "Po1" "Po2" "LF" "M.F"
## [8] "Pop" "NW" "U1" "U2" "Wealth" "Ineq" "Prob"
## [15] "Time" "Crime"
```

```
# In the data, the crime column is the response as Y, all the rest columns are variable X
# To build the model use LM
model_lm <- lm(Crime ~ M+So+Ed+Po1+Po2+LF+M.F+NW+U1+U2+Wealth+Ineq+Prob+Time, data=df_crime)

model_lm
```

```
##
## Call:
## lm(formula = Crime ~ M + So + Ed + Po1 + Po2 + LF + M.F + NW +
## U1 + U2 + Wealth + Ineq + Prob + Time, data = df_crime)
##
## Coefficients:
## (Intercept) M So Ed Po1
## -6.216e+03 8.978e+01 -4.105e+00 1.888e+02 1.896e+02
## Po2 LF M.F NW U1
## -1.122e+02 -7.845e+02 2.188e+01 4.335e+00 -6.253e+03
## U2 Wealth Ineq Prob Time
## 1.695e+02 9.103e-02 6.745e+01 -4.858e+03 -4.456e+00
```

```
# evaluate model lm
summary(model_lm)
```

```
##
## Call:
## lm(formula = Crime ~ M + So + Ed + Po1 + Po2 + LF + M.F + NW +
##      U1 + U2 + Wealth + Ineq + Prob + Time, data = df_crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -380.91 -101.89  -14.77   110.87   505.40
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.216e+03  1.560e+03  -3.986 0.000364 ***
## M              8.978e+01  4.113e+01   2.183 0.036491 *
## So            -4.105e+00  1.472e+02  -0.028 0.977921
## Ed             1.888e+02  6.142e+01   3.074 0.004293 **
## Po1            1.896e+02  1.048e+02   1.808 0.079978 .
## Po2           -1.122e+02  1.161e+02  -0.966 0.341042
## LF            -7.845e+02  1.439e+03  -0.545 0.589392
## M.F            2.188e+01  1.857e+01   1.178 0.247367
## NW             4.335e+00  6.408e+00   0.676 0.503623
## U1            -6.253e+03  4.099e+03  -1.525 0.136993
## U2             1.695e+02  8.140e+01   2.083 0.045351 *
## Wealth         9.103e-02  1.022e-01   0.891 0.379612
## Ineq           6.745e+01  2.176e+01   3.099 0.004026 **
## Prob          -4.858e+03  2.248e+03  -2.161 0.038299 *
## Time          -4.456e+00  6.882e+00  -0.648 0.521894
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 206.8 on 32 degrees of freedom
## Multiple R-squared:  0.801, Adjusted R-squared:  0.714
## F-statistic: 9.202 on 14 and 32 DF, p-value: 1.301e-07
```

```
# Fit the input data into model
```

```
# Create predict input as given
```

```
predict_input <- data.frame(M = 14.0, So = 0, Ed = 10.0, Po1 = 12.0, Po2 = 15.5, LF = 0.640, M.F = 94.0, P
```

```
# Fit into lm model
```

```
predict_crime <- predict(model_lm, predict_input)
```

```
predict_crime
```

```
##      1
## 162.4041
```

From above, the predicted crime is 162.4041.


To use the solution of 5.1 to evaluate if the predicted value is outlier because the value seems too low

First to add the new value to the existing crime data to draw boxplot

```
library(ggplot2)
```

```
## Registered S3 methods overwritten by 'ggplot2':  
##   method      from  
## [.quosures    rlang  
## c.quosures    rlang  
## print.quosures rlang
```

```
df_crime_new <-c(df_crime[,16], 162.4041)  
qplot(y=df_crime_new, x= 1, geom = "boxplot")
```



Then use Grubbs Test to find the outliers

```
# use type = 10 to test for one outlier  
library(outliers)  
test_11 <- grubbs.test(df_crime_new, type = 11)  
test_11
```

```
##  
## Grubbs test for two opposite outliers  
##  
## data: df_crime_new  
## G = 4.60691, U = 0.76427, p-value = 0.7254  
## alternative hypothesis: 162.4041 and 1993 are outliers
```

From the above test, it shows the predicted value has a good chance to be as an outlier because it is too low.

Update the *lm* model, noticed from the summary that *p* value of different factors are all different. Choose the factors with small *p*-value: *M*, *So*, *Ed*, *Po1*, *U2*, *Ineq*, *Prob*

```
model_lm_2 <- lm(Crime ~ M+So+Ed+Po1+U2+Ineq+Prob, data=df_crime)
model_lm_2
```

```
##
## Call:
## lm(formula = Crime ~ M + So + Ed + Po1 + U2 + Ineq + Prob, data = df_crime)
##
## Coefficients:
## (Intercept)          M          So          Ed          Po1
##    -4959.30      99.51      73.01     205.24     111.96
##          U2        Ineq        Prob
##      88.48      63.87    -4223.04
```

```
summary(model_lm_2)
```

```
##
## Call:
## lm(formula = Crime ~ M + So + Ed + Po1 + U2 + Ineq + Prob, data = df_crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -480.55  -79.12  -14.81   122.64   562.66
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4959.30      914.45  -5.423 3.27e-06 ***
## M              99.51       34.55   2.880 0.006422 **
## So             73.01      110.11   0.663 0.511200
## Ed            205.24       46.97   4.370 8.94e-05 ***
## Po1           111.96       14.60   7.666 2.65e-09 ***
## U2             88.48       41.22   2.147 0.038105 *
## Ineq          63.87       15.15   4.216 0.000143 ***
## Prob        -4223.04     1664.87  -2.537 0.015310 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 202.1 on 39 degrees of freedom
## Multiple R-squared:  0.7685, Adjusted R-squared:  0.7269
## F-statistic: 18.49 on 7 and 39 DF,  p-value: 1.368e-10
```

Use the new model to do the prediction

```
predict_crime_2 <- predict(model_lm_2, predict_input)
predict_crime_2
```

```
##          1
## 1263.083
```

The new data seems a better range. Double check on the outlier

```
df_crime_new <-c(df_crime[,16], 1263.083)
test_2 <- grubbs.test(df_crime_new, type = 11)
test_2
```

```
##
## Grubbs test for two opposite outliers
##
## data: df_crime_new
## G = 4.27610, U = 0.78612, p-value = 1
## alternative hypothesis: 342 and 1993 are outliers
```

From above, the new predicted value is not an outlier.

*Show model output and quality of fit

```
summary(model_lm_2)
```

```
##
## Call:
## lm(formula = Crime ~ M + So + Ed + Po1 + U2 + Ineq + Prob, data = df_crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -480.55  -79.12  -14.81   122.64   562.66
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4959.30     914.45  -5.423 3.27e-06 ***
## M              99.51      34.55   2.880 0.006422 **
## So             73.01     110.11   0.663 0.511200
## Ed            205.24      46.97   4.370 8.94e-05 ***
## Po1           111.96      14.60   7.666 2.65e-09 ***
## U2             88.48      41.22   2.147 0.038105 *
## Ineq           63.87      15.15   4.216 0.000143 ***
## Prob        -4223.04    1664.87  -2.537 0.015310 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 202.1 on 39 degrees of freedom
## Multiple R-squared:  0.7685, Adjusted R-squared:  0.7269
## F-statistic: 18.49 on 7 and 39 DF,  p-value: 1.368e-10
```

The predicted crime is 1263.083

The R-squared is over 0.72 which means the linear model can explain over 72% of the prediction.